


Architecture Overview - Traefik + Nginx Explained

System Architecture



How It Works

Request Flow

1. User visits <http://localhost/>

Browser → Traefik → Frontend (Nginx) → Returns index.html + React app

2. User clicks "Login" (stays in React app)

React Router handles /login → No server request

3. User submits login form

React → Axios POST http://localhost/api/auth/login
→ Traefik sees /api → Routes to Backend
→ Backend processes → Returns JWT token

4. User creates project

React → Axios POST http://localhost/api/projects
→ Traefik sees /api → Routes to Backend
→ Backend → PostgreSQL → Returns data

Component Roles

Traefik (Reverse Proxy)

Purpose: Smart router that directs traffic

What it does:

- Listens on port 80
- Reads Docker labels from containers
- Routes `/api/*` → Backend container
- Routes everything else → Frontend container
- No configuration files needed (uses labels)

Why we use it:

- Automatic service discovery
- Zero-downtime deployments
- Load balancing (if multiple instances)
- Easy HTTPS with Let's Encrypt

Nginx (Static File Server)

Purpose: Serves React build files

What it does:

- Runs inside frontend container
- Serves `/usr/share/nginx/html` (React build output)
- Handles SPA routing (all routes → index.html)
- Gzip compression
- Static file caching

Why we use it:

- Fast static file serving
- Production-ready
- Low resource usage
- Battle-tested

Backend (Node.js + Express)

Purpose: REST API server

What it does:

- Handles all `/api/*` requests
- Authenticates users (JWT)
- Manages database operations
- Calls AI provider APIs
- Returns JSON responses

Frontend (React + Vite)

Purpose: User interface

What it does:

- Single Page Application (SPA)
- Client-side routing (React Router)
- Makes API calls to backend
- Renders UI components
- Manages local state

Port Mapping

Service	Internal Port	External Port	Access
Traefik	80	80	http://localhost
Traefik Dashboard	8080	8080	http://localhost:8080
Backend	3000	-	Via Traefik only
Frontend (Nginx)	80	-	Via Traefik only
PostgreSQL	5432	-	Internal only

Key Point: Only Traefik is exposed. Everything else is internal to Docker network.

Routing Rules (Traefik Labels)

Backend Labels

```
yaml
labels:
  - "traefik.enable=true"
  - "traefik.http.routers.backend.rule=Host('localhost') && PathPrefix('/api')"
  - "traefik.http.services.backend.loadbalancer.server.port=3000"
```

Meaning:

- Enable Traefik routing for this container
- Match requests to `localhost` with path starting with `/api`
- Forward to container's port 3000

Frontend Labels

```
yaml
labels:
  - "traefik.enable=true"
  - "traefik.http.routers.frontend.rule=Host('localhost')"
  - "traefik.http.services.frontend.loadbalancer.server.port=80"
```

Meaning:

- Enable Traefik routing for this container
- Match all requests to `localhost` (catch-all)
- Forward to container's port 80 (Nginx)
- Lower priority than backend (PathPrefix is more specific)

URL Examples

Frontend URLs (Handled by React Router)

```
http://localhost/      → Frontend (index.html)
http://localhost/login → Frontend (index.html)
http://localhost/dashboard → Frontend (index.html)
http://localhost/project/1 → Frontend (index.html)
http://localhost/admin   → Frontend (index.html)
```

All these return the same `(index.html)`, then React Router takes over.

Backend API URLs

```
http://localhost/api/health → Backend
http://localhost/api/auth/login → Backend
http://localhost/api/projects → Backend
http://localhost/api/chats/1/messages → Backend
http://localhost/api/providers/keys → Backend
```

All `(/api/*)` requests go to the backend.

Development vs Production

Development Mode

Frontend:

```
bash
cd frontend
npm run dev # Vite dev server on port 5173
```

- Hot module replacement
- Source maps
- Fast refresh

Backend:

```
bash
cd backend
npm run dev # Nodemon on port 3000
```

- Auto-restart on changes
- Direct database connection

Access:

- Frontend: <http://localhost:5173>
- Backend: <http://localhost:3000>
- (No Traefik in development)

Production Mode (Docker)

```
bash  
docker-compose up -d
```

Frontend:

- Built with Vite (`(npm run build)`)
- Static files served by Nginx
- Gzipped, cached, optimized

Backend:

- Runs with `(npm start)`
- Production mode (`(NODE_ENV=production)`)

Access:

- Everything: <http://localhost> (via Traefik)

🚀 Scaling Options

Multiple Backend Instances

```
yaml  
  
backend:  
  deploy:  
    replicas: 3
```

Traefik automatically load balances between them.

Multiple Frontend Instances

Not needed - Nginx serves static files efficiently.

Database Replication

Can add read replicas for PostgreSQL if needed.

Security Layers

1. Network Isolation

All containers on private `ai-platform` network. Only Traefik exposed to host.

2. Traefik (Edge)

- Can add rate limiting
- Can add IP whitelist
- Can add authentication middleware

3. Backend (API)

- JWT authentication
- Admin role checks
- Input validation
- SQL injection prevention (parameterized queries)

4. Database

- Not exposed to host
- Only backend can connect
- User permissions

Nginx Configuration Explained

```
nginx
```

```
server {
    listen 80;
    root /usr/share/nginx/html;

    # Try to serve file, fallback to index.html for SPA routing
    location / {
        try_files $uri $uri/ /index.html;
    }

    # Cache static assets
    location ~* \.(js|css|png|jpg|svg)$ {
        expires 1y;
    }
}
```

Why `try_files $uri $uri/ /index.html`?

React Router uses client-side routing. When user visits `/dashboard`:

1. Nginx checks if file `/dashboard` exists → No
2. Nginx checks if directory `/dashboard/` exists → No
3. Nginx serves `/index.html` → React app loads → React Router shows Dashboard

Without this, refreshing on `/dashboard` would give 404.

💡 Testing the Architecture

Test Traefik Routing

```
bash

# Should return React app HTML
curl http://localhost/

# Should return API JSON
curl http://localhost/api/health

# Traefik dashboard
open http://localhost:8080
```

Test Backend Directly (Development)

```
bash
```

```
# If backend running outside Docker  
curl http://localhost:3000/api/health
```

Test Frontend Directly (Development)

```
bash
```

```
# If frontend running with Vite  
open http://localhost:5173
```

Common Issues

"Cannot reach backend from frontend"

Problem: Frontend tries to call `(http://backend:3000/api)`

Solution: Use relative URLs `(/api)` not absolute URLs. Traefik handles routing.

```
javascript
```

```
// ✗ Wrong  
axios.get('http://backend:3000/api/health')  
  
// ✓ Correct  
axios.get('/api/health')
```

"404 on page refresh"

Problem: Nginx not configured for SPA routing

Solution: Check nginx.conf has `(try_files $uri $uri/ /index.html;)`

"Traefik shows no routes"

Problem: Labels not applied

Solution:

```
bash
```

```
# Check Traefik dashboard  
open http://localhost:8080  
  
# Restart services  
docker-compose restart
```

Component	Technology	Purpose	Port
Traefik	Go	Reverse proxy & router	80
Nginx	C	Static file server	80 (internal)
Backend	Node.js	REST API	3000 (internal)
Frontend	React	UI	Built files
PostgreSQL	SQL	Database	5432 (internal)

Key Takeaway:

- Traefik = Smart router (reads Docker labels)
- Nginx = Fast file server (inside frontend container)
- They work together, not in competition!