


How to Add a New AI Provider

The system now supports **dynamic AI provider addition** with a plugin architecture. You can add new providers without modifying core code!

Architecture Overview



Key Components:

1. **BaseAdapter** - Abstract class defining the interface all providers must implement
2. **Provider Adapters** - Individual files for each AI provider (openai.js, anthropic.js, etc.)
3. **aiProviders Service** - Dynamically loads adapters based on database configuration
4. **Database** - Stores provider configuration including which adapter to use

Step-by-Step: Adding a New Provider

Let's add **Mistral AI** as an example.

Step 1: Create the Adapter File

Create `backend/src/services/adapters/mistral.js`:

```
javascript
```

```
const BaseAdapter = require('./BaseAdapter');

class MistralAdapter extends BaseAdapter {
  // Transform messages to Mistral format
  transformMessages(messages) {
    return messages.map(m => ({
      role: m.role,
      content: m.content
    }));
  }

  // Build the API request payload
  buildRequestPayload(messages) {
    return {
      model: this.modelName,
      messages: this.transformMessages(messages),
      temperature: this.config.temperature || 0.7,
      max_tokens: this.config.max_tokens || 1000
    };
  }
}

// Set HTTP headers for authentication
getHeaders(apiKey) {
  return {
    'Authorization': `Bearer ${apiKey}`,
    'Content-Type': 'application/json'
  };
}

// Extract the AI response text
extractResponse(responseData) {
  return responseData.choices[0].message.content;
}

// Extract token usage (optional)
extractTokenCount(responseData) {
  return responseData.usage?.total_tokens || 0;
}

module.exports = MistralAdapter;
```

Step 2: Add Provider to Database

sql

```
INSERT INTO ai_providers (
    name,
    display_name,
    api_endpoint,
    model_name,
    adapter_module,
    config,
    is_active
) VALUES (
    'mistral', -- unique identifier
    'Mistral AI', -- display name
    'https://api.mistral.ai/v1/chat/completions', -- API endpoint
    'mistral-medium', -- default model
    'mistral', -- adapter file name (without .js)
    '{"temperature": 0.7, "max_tokens": 1000}', -- JSON configuration
    true -- is active
);
```

Or use the API:

bash

```
curl -X POST http://localhost/api/providers/admin/create \
-H "Authorization: Bearer YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '{
    "name": "mistral",
    "display_name": "Mistral AI",
    "api_endpoint": "https://api.mistral.ai/v1/chat/completions",
    "model_name": "mistral-medium",
    "adapter_module": "mistral",
    "config": {
        "temperature": 0.7,
        "max_tokens": 1000
    }
}'
```

Step 3: Add API Keys

```
bash
```

```
# Get the provider ID (let's say it's 5)
curl http://localhost/api/providers

# Add API key
curl -X POST http://localhost/api/providers/keys \
-H "Authorization: Bearer YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "provider_id": 5,
  "key_value": "your-mistral-api-key"
}'
```

Step 4: Test It!

```
bash
```

```
# Create a chat with Mistral
curl -X POST http://localhost/api/chats/project/1 \
-H "Authorization: Bearer YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "title": "Testing Mistral",
  "provider_id": 5
}'

# Send a message
curl -X POST http://localhost/api/chats/CHAT_ID/messages \
-H "Authorization: Bearer YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "content": "Hello Mistral!"
}'
```

BaseAdapter Methods Reference

Every adapter **must implement** these methods:

Required Methods

```
javascript
```

```
transformMessages(messages)
// Input: [{ role: 'user', content: '...' }, ...]
// Output: Provider-specific message format
// Purpose: Convert our standard format to provider's format

buildRequestPayload(messages)
// Input: Array of messages
// Output: Complete request body object
// Purpose: Build the full API request payload

getHeaders(apiKey)
// Input: API key string
// Output: Object with HTTP headers
// Purpose: Set authentication and content-type headers

extractResponse(responseData)
// Input: Full API response data
// Output: String (the AI's response text)
// Purpose: Extract the actual message content
```

Optional Methods

```
javascript
```

```
extractTokenCount(responseData)
// Input: Full API response data
// Output: Number (token count)
// Default: Returns 0
// Purpose: Extract token usage for analytics

getEndpoint(apiKey)
// Input: API key string
// Output: String (full URL)
// Default: Uses this.endpoint from database
// Purpose: Customize endpoint URL (e.g., add API key as query param)

parseRateLimitError(error)
// Input: Axios error object
// Output: { resetAt: Date } or null
// Default: BaseAdapter provides standard implementation
// Purpose: Extract rate limit reset time from error
```

Real-World Examples

Example 1: OpenAI-Compatible Provider (Groq, Together.ai)

Most providers following OpenAI's API format need minimal code:

javascript

```
const BaseAdapter = require('./BaseAdapter');

class GroqAdapter extends BaseAdapter {
  transformMessages(messages) {
    return messages.map(m => ({ role: m.role, content: m.content }));
  }

  buildRequestPayload(messages) {
    return {
      model: this.modelName,
      messages: this.transformMessages(messages)
    };
  }

  getHeaders(apiKey) {
    return {
      'Authorization': `Bearer ${apiKey}`,
      'Content-Type': 'application/json'
    };
  }

  extractResponse(responseData) {
    return responseData.choices[0].message.content;
  }

  extractTokenCount(responseData) {
    return responseData.usage?.total_tokens || 0;
  }
}

module.exports = GroqAdapter;
```

Example 2: Provider with API Key in URL (Like Gemini)

```
javascript
```

```
class CustomAdapter extends BaseAdapter {  
    // ... other methods ...  
  
    getEndpoint(apiKey) {  
        // Add API key to URL instead of header  
        return `${this.endpoint}?apiKey=${apiKey}&format=json`;  
    }  
  
    getHeaders(apiKey) {  
        // No auth header needed  
        return {  
            'Content-Type': 'application/json'  
        };  
    }  
}
```

Example 3: Provider with Custom Message Format

```
javascript
```

```
class CustomAdapter extends BaseAdapter {  
    transformMessages(messages) {  
        // Custom format with nested structure  
        return {  
            conversation: messages.map((m, idx) => ({  
                id: idx,  
                speaker: m.role === 'user' ? 'human' : 'ai',  
                text: m.content,  
                timestamp: new Date().toISOString()  
            }))  
        };  
    }  
  
    buildRequestPayload(messages) {  
        return {  
            model_id: this.modelName,  
            input: this.transformMessages(messages),  
            parameters: this.config  
        };  
    }  
  
    extractResponse(responseData) {  
        // Navigate nested response structure  
        return responseData.output.generated_text.content;  
    }  
}
```

Configuration Options

The `config` JSONB field in `(ai_providers)` table allows provider-specific settings:

```
javascript

{
  "temperature": 0.7,      // Control randomness
  "max_tokens": 1000,     // Maximum response length
  "top_p": 0.9,          // Nucleus sampling
  "frequency_penalty": 0, // Reduce repetition
  "presence_penalty": 0,  // Encourage new topics
  "custom_param": "value" // Any provider-specific parameter
}
```

Access in adapter via `(this.config)`:

```
javascript

buildRequestPayload(messages) {
  return {
    model: this.modelName,
    messages: this.transformMessages(messages),
    temperature: this.config.temperature || 0.7,
    my_custom_setting: this.config.custom_param
  };
}
```

Testing Your Adapter

1. Test Locally First

javascript

```
// Create a test file: backend/test-adapter.js
const MistralAdapter = require('./src/services/adapters/mistral');

const mockProvider = {
  name: 'mistral',
  api_endpoint: 'https://api.mistral.ai/v1/chat/completions',
  model_name: 'mistral-medium',
  config: { temperature: 0.7 }
};

const adapter = new MistralAdapter(mockProvider);

const messages = [
  { role: 'user', content: 'Hello!' }
];

console.log('Transformed:', adapter.transformMessages(messages));
console.log('Payload:', adapter.buildRequestPayload(messages));
console.log('Headers:', adapter.getHeaders('test-key'));
```

Run: `node backend/test-adapter.js`

2. Test with Real API

Use the provider management endpoints to add it to the system and test end-to-end.

Troubleshooting

Adapter Not Found

Error: Adapter module 'mistral' not found

Solution: Ensure the file exists at `backend/src/services/adapters/mistral.js`

Method Not Implemented

Error: transformMessages() must be implemented by adapter

Solution: Make sure all required methods are implemented in your adapter

Rate Limit Handling

If the provider has custom rate limit headers, override `(parseRateLimitError())`:

```
javascript
```

```
parseRateLimitError(error) {
  if (error.response?.status !== 429) return null;

  const resetTime = error.response.headers['x-custom-reset-header'];
  const resetAt = new Date(resetTime * 1000);

  return { resetAt };
}
```

Benefits of This Architecture

- ✓ No Core Code Changes** - Add providers without touching main codebase
- ✓ Consistent Interface** - All providers work the same way from the app's perspective
- ✓ Easy Maintenance** - Each provider's logic is isolated
- ✓ Type Safety** - BaseAdapter enforces required methods
- ✓ Flexible Configuration** - Store provider-specific settings in database
- ✓ Hot Reload Ready** - Can add providers without restarting (in development)

Advanced: Dynamic Model Selection

You can even support multiple models per provider:

```
sql
-- Add multiple entries for same provider with different models
INSERT INTO ai_providers (name, display_name, model_name, adapter_module) VALUES
('gpt-4', 'GPT-4', 'gpt-4', 'openai'),
('gpt-3.5', 'GPT-3.5 Turbo', 'gpt-3.5-turbo', 'openai');
```

Users can then choose between GPT-4 and GPT-3.5 as separate options!